

The Complete Auditor Flow

Schedule to Verified Closure

The story of one audit, from the calendar that booked it to the day its last corrective action was verified on site.

Sources	EHSERA-ISMS-AO-2026-v1.0 · ALGO-ISMS-WF-2026-v1.0 · AUD-FORM-01
Standards	ISO 45001 Cl.9 · ISO 14001 Cl.9 · OSHA VPP
Surfaces	Mobile (React Native) — the audit. Web console (React) — before and after.
Roles	Auditor · Worker · Supervisor · Safety Manager (mobile + web) · Admin (web only) · The System
Worked example	AUD-000010 — Q3 Fire Safety Audit, WindTech Nacelle Manufacturing Ltd

Every number in this document is real.

AUD-000010 was run end to end through the built system against the live database. The score, the classifications, the escalations that fired and the timestamps are what the software actually produced — not an illustration of what it might produce.

Before the story starts

What kind of thing an audit is, in this system.

An audit here is not a form somebody fills in. It is a ten-step process with named owners, hard stops that cannot be skipped, and a set of rules that fire on their own without anyone having to notice. Four things follow from that, and they shape everything below.

Nobody books an audit by hand

"Audits are not booked by hand. The system generates the annual programme from each site's risk band, and that band is driven by the site's own safety performance score. A site that deteriorates gets audited more often, automatically."

The calendar is generated from data the platform already holds. A site whose safety performance score climbs into the critical band starts receiving monthly inspections and quarterly audits without a person deciding it should.

The audit happens on the phone, in the field, as it happens

"Every observation logged live on the app — no paper, no writing up afterwards."

Steps 4 to 8 run on the phone while the auditor is standing in the place they are auditing. There is no write-up stage afterwards, because a write-up stage is where detail goes to die and where the gap between what was seen and what was recorded opens up.

The auditor judges; the system does the arithmetic

The auditor decides what each finding is. The system calculates the score, applies the thresholds, flags repeats, and raises the corrective actions. Neither does the other's job — and two rules are taken out of the auditor's hands entirely, which we come to at step 7.

Issuing the report is not the end

"An audit is not closed when the report is issued — it stays open until every corrective action it raised has been verified effective."

This is the sentence that most audit software gets wrong. Marking an action complete is a claim made by whoever owned it. Verification is somebody standing in the same place afterwards, checking the fix is holding. The two are different events and this system keeps them apart.

STEP 01 The calendar builds itself

PLAN · The System · automatic — no human involved

The story begins with nobody. On the audit programme screen, the platform has already scored every site in the organisation and sorted them into four risk bands. WindTech Nacelle Manufacturing Ltd sits at a safety performance score of 43.1, which places it in the medium band: quarterly inspection, bi-annual audit.

The band is not a label somebody typed. It comes from the same KPI engine the safety dashboards read, so a site cannot be critical on one screen and low on another.

Band	Qualifying criteria	How often	Re-audit trigger
CRITICAL	Score 75+ or any fatal/critical event in 12 months	Monthly inspection + quarterly full audit	Any Major NC → re-audit within 30 days
HIGH	Score 50-74 or a lost-time injury in 6 months	Monthly inspection + quarterly audit	Two or more Major NCs in 12 months → 60 days
MEDIUM	Score 25-49	Quarterly inspection + bi-annual audit	Overall score below 65% for two consecutive audits
LOW	Score 0-24, no serious event in 24 months	Bi-annual inspection + annual audit	Any lost-time injury or score below 65% → upgraded to Medium

Computed is not the same as authorised

The system works out the cadence. It does not act on it until a person signs it off. The Safety Manager authorises the programme for their site — optionally attaching a specific concern to be included in scope, which travels all the way to the auditor's brief pack. The Admin then approves the calendar across all sites.

Until that authorisation exists, pressing Generate produces nothing and says why. In testing, all nine sites refused generation with “The programme for this site has not been authorised”. Generating a year of work nobody signed off is the booking-by-hand problem wearing a different hat.

Once authorised, generation produces the year:

Site	Band	Generated for 2027	Re-run
WindTech Nacelle Manufacturing Ltd	MEDIUM	5 inspections + 3 audits = 8 events	0 created, 8 skipped

Generation is idempotent, deliberately.

It counts what is already booked in each window and fills only the gap. Running it twice in March produces the same calendar as running it once — and it never touches an audit that has already started, because an audit already walked is a record, not a plan.

WEB

Audit Programme — risk band table, per-site authorisation, Generate, Admin approval, and the 14-day reminder sweep.

STEP 02 Somebody has to be independent

PLAN · Safety Manager · HARD STOP

AUD-000010 now exists in the register with no auditor against it, and it cannot move. The Safety Manager names who audits what — and this is the step the specification marks in red.

“Names who audits what, and must ensure the auditor is independent of the area being audited.”

This is what makes the finding credible.”

That is enforced, not advised. Assigning an auditor who is the supervisor of the area under audit is refused outright:

“The auditor cannot be the supervisor of the area being audited — independence is what makes the finding credible.” — the system's actual refusal, HTTP 400

The moment a valid auditor is named, two notifications go out. The auditor learns the audit is theirs and that their brief pack arrives seven days before the visit. The supervisor of the area — the auditee — is told an audit is coming, with at least two weeks' notice.

Unless it is meant to be a surprise.

Two weeks' notice applies to scheduled, management-directed and regulatory audits. It does not apply to an inspection raised by a risk spike, a score threshold breach or an incident — those carry no notice by design, and the system stamps nothing rather than recording a notice that was never given.

Six things can start an audit, and which one it was changes how it behaves:

Trigger	What it means	Notice?
Scheduled programme	The annual calendar generated from the site risk band	Yes
After an incident	A reactive inspection following a significant event	No
Management directed	Requested by management outside the normal cycle	Yes
Regulatory requirement	Required by an authority or for certification renewal	Yes
Score threshold breach	The safety performance score crosses a threshold	No
Risk spike	Predicted risk rises sharply — raised rather than waiting	No

WEB

Audit Register — the queue, the step each audit waits on, and the assign dialog showing each auditor's workload and qualifications.

STEP 03 The auditor is briefed before they leave

PREPARE · Auditor · on the phone, readable offline

Seven days before the visit, the system builds the brief pack. The auditor does not assemble it and cannot edit it. It arrives in their app already containing everything they should walk in knowing.

For AUD-000010 the pack contained:

- previous findings from the last two audits at this site, and whether each was closed
- open corrective actions still outstanding, flagged where overdue
- the current score — what this site achieved last time, and its rating
- 25 permits past their validity date at this site, drawn live from the permit register
- the areas that failed hardest last time, used to reorder the checklist
- the ISO clauses this audit type maps to

It is a snapshot, on purpose.

The pack is generated once and stored, rather than recomputed each time it is opened. The record has to show what the auditor was actually briefed on — not what the data happens to look like when somebody opens it months later.

The checklist leads with last time's failures

Sections that produced non-conformances in the previous two audits are weighted, and the checklist is ordered to put them first. An auditor who runs out of time runs out of it at the bottom of the list, so what sits at the top matters.

A closed finding is a watchlist item

Findings that were signed off last time appear in a separate repeat watchlist. If the auditor finds one again, the system flags it as a repeat and treats it as more serious than a first occurrence — the control was already supposed to be in place.

The only action on this screen is confirming the brief has been read. That confirmation is what unlocks the opening meeting: an auditor who has not seen last time's findings cannot set a credible scope.

MOBILE Brief Pack — the whole pack, readable with no signal, ending in “I have read the brief”.

STEP 04 Scope is agreed in the room, not afterwards

CONDUCT · Auditor + Supervisor · on site

The auditor arrives. Before anything is inspected, they hold the opening meeting with the supervisor of the area, and the app captures it as a structured record rather than a note.

Three fields are required, and the app refuses to continue without them:

Field	AUD-000010's actual entry
Scope	Assembly hall fire detection, egress and extinguishers. Paint shop out of scope.
Method	Physical walk, worker interview, record sampling
Sampling approach	All 12 extinguishers; 3 of 9 exit routes at random; 5 worker interviews

They are three separate required fields rather than one notes box because the entire value of the opening meeting is that there is no dispute afterwards about what was in or out of scope — and free text does not settle that argument.

The supervisor attends, and their attendance is recorded either way. If they are not there the audit is not blocked, but the fact is stored: “scope was agreed jointly” is a claim the report makes, and it has to be true. The meeting is GPS-stamped.

MOBILE

Opening Meeting — scope, method, sampling, attendees, and whether the supervisor was present.

STEPS 05-06 The walk

CONDUCT · Auditor · one item per screen, offline-capable

Now the auditor walks the site. The app shows one checklist item per screen and they swipe through. Each item carries its section, its ISO clause, and — where it applies — a red flag marking it as critical.

AUD-000010's fire safety checklist had four items. Here is what happened, in order:

#	Item	Answer	Running score
1	Extinguishers (CRITICAL)	Full — 2 pts	100% Excellent
2	Exit Routes (CRITICAL)	None — 0 pts	50% Poor
3	Alarm System	Partial — 1 pt	50% Poor
4	Emergency Lighting	Full — 2 pts	62.5% Acceptable

What happened when item 2 scored zero

Exit Route 3 was blocked by stacked pallets. The auditor tapped None. Because that item is marked critical, three things happened before their thumb left the screen:

- the audit moved into stop-work, which the eight-stage engine reads as RESPOND
- the Safety Manager and the executive were notified immediately — not at submit, not in the report
- the app showed the auditor exactly what had been raised and by what rule

“Exit Routes’ is a critical item and scored zero. The Safety Manager and the executive have been notified. Work may be suspended before this audit finishes.” — what the app actually displayed

The auditor confirms the hazard is contained and the walk resumes — it does not restart. Containment happens first; the audit carries on afterwards.

The scoring rubric

Answers are not pass/fail. A full conformance is worth twice an observation, and items marked Not Applicable leave the calculation entirely so a score is never diluted by questions that did not apply to that site.

Answer	Points	Meaning
Full compliance	2	Meets the requirement
Partial / observation	1	Compliant but improvable
Non-compliance	0	Does not meet the requirement
Not applicable	excluded	Removed from the denominator entirely

Both specifications express the formula differently and they are mathematically identical. The app ships both so an auditor holding either document recognises what the phone is showing:

Specification	Formula
Algorithmic spec (ALGO-ISMS-WF-2026)	$(\text{points earned} / \text{points possible}) \times 100$
Form specification (AUD-FORM-01)	$(C + 0.5 \times \text{OBS}) / \text{total assessed} \times 100$

Evidence is attached to the line it proves

“Photos attached to the specific checklist item, not dumped in a general folder.”

On AUD-000010 the auditor attached three pieces of evidence, each bound to its checklist line:

Kind	Attached to	What it recorded
------	-------------	------------------

Note	Exit Routes	Pallets stacked across exit 3
Scan	Alarm System	Asset reference AST-EXT-0442
Interview	Exit Routes	Worker named exit 3 — the blocked one. Did not know the alternate. Competence card not verified.

A worker interview is evidence, not a note.

The specification is blunt about this: what the worker actually does is the evidence, not what the procedure says. So the answer is recorded against the person who gave it, against the checklist line it proves, with the prompt they were asked and whether their competence card checked out against the matrix.

Every answer and every piece of evidence is GPS-stamped, which is what makes the observation defensible to a regulator: it proves the observation was made at the place claimed.

The whole walk works with no signal.

Audits happen in plant areas, tank farms and remote sites. Answers are held on the device the instant they are tapped. Where there is signal they also go up immediately — that is what returns the running score and fires the critical alert while work can still be suspended. Where there is not, the answer is marked unsent and the whole set is carried up at step 7.

MOBILE

Field Inspection — one item per screen, swipe navigation, running score in the header, evidence sheet, and the critical-item alert.

STEP 07 The auditor judges; the system counts

CLASSIFY · Auditor · HARD STOP

With the walk finished, the auditor reviews every answer and assigns a classification to each finding. The app suggests one; the auditor decides.

Classification	What it means	Action due
CONFORMANCE	Meets the requirement. Recorded as a positive — audits record what is working, not only what is wrong.	—
OBSERVATION	Compliant but improvable. A finding, but not a non-conformance.	—
MINOR NC	A lapse that does not undermine the system.	30 days
MAJOR NC	A systemic failure. Safety Manager notified within 24 hours.	7 days
CRITICAL / REGULATORY	Immediate danger or a legal breach. Executive notified at once and work may be suspended.	1 day

Two judgements are not the auditor's to make

The auditor owns every classification except where a rule overrides them, and the app says so rather than silently correcting:

- **Critical zeros.** A critical item scoring zero is a Major NC at minimum. It can be escalated to Critical; it cannot be softened below Major. That is the entire point of marking an item critical.
- **Weak sections.** A section scoring below 60% raises a Minor NC of its own, attributed to nobody and impossible to suppress. A section falling below the threshold is a lapse in the system, not in one item.

On AUD-000010, both rules fired. The Egress section scored 50%:

Section	Score	Outcome
Detection & Suppression	75%	—
Egress	50%	Below the 60% threshold → automatic Minor NC

The five findings the system produced:

Ref	Finding	Classification	Action due
AUD-000010-F01	Extinguishers	Conformance	—
AUD-000010-F02	Exit Routes	MAJOR NC (critical item, auto)	28 Aug — 7 days
AUD-000010-F03	Alarm System	Observation	—
AUD-000010-F04	Emergency Lighting	Conformance	—
AUD-000010-F05	Section 'Egress' scored 50%	MINOR NC (section rule, auto)	20 Sep — 30 days

The score and the rating are different things

The system computed 62.5% — five points earned out of eight possible — which is the Acceptable band. The overall rating came back Unsatisfactory.

62% and Unsatisfactory at the same time.

That is not a contradiction; it is the design. The rating is set from the finding counts, not from the percentage. Any Major non-conformance or regulatory breach makes an audit unsatisfactory whatever it scored — because a good average with a systemic failure inside it is still a site with a systemic failure. More than three Minor NCs makes it Requires Improvement.

Band	Range	AUD-000010
EXCELLENT	90% and above	
GOOD	75 – 89%	
ACCEPTABLE	60 – 74%	62.5% ←
POOR	Below 60%	

Two thresholds sit alongside the bands and act on their own: below 70% automatically alerts the Safety Manager, and below 65% twice in a row at one site forces a re-audit. At 62%, the first of those fired.

MOBILE

Findings & Score — the arithmetic shown but not editable, the classification chips that are, and the two rules the auditor cannot override.

STEP 08 The auditee gets to say it is wrong

AGREE · Auditor + Supervisor · HARD STOP · before anyone leaves site

The closing meeting is where the supervisor sees the findings — on the auditor's screen, in front of them — and gets their one opportunity to correct a factual error before anything is fixed or reported.

"Their opportunity to correct a factual error before anything is fixed. After this meeting findings are locked and can only change through a formal amendment."

Three things have to happen, and the app will not finish without all three:

- **Presented.** every finding is presented on screen to the supervisor
- **Confirmed.** the supervisor confirms the findings are factually accurate — or disputes them
- **Signed.** both parties sign on the device, and the findings lock immediately

Disputing is a real outcome

If the supervisor says a finding is factually wrong, the dispute is recorded and nothing locks. The auditor goes back, corrects the error and holds the meeting again. That was tested on AUD-000010: a disputed closing meeting left findings unlocked and the audit sitting at step 8, exactly as it should.

Both signatures, on the device, before anyone leaves

Signing later — on paper, or by email the next day — is the gap through which findings drift between the walk and the report. So the signature pad is on the phone, and capturing both signatures is what locks the findings.

The corrective action deadlines are also agreed here, in front of the supervisor, and they override the system's defaults. The supervisor committed to those dates with the auditor standing there.

After the meeting, the checklist is frozen. Attempting to change an answer returns:

"Findings were locked at the closing meeting and can only change through a formal amendment."
— the system's actual refusal, HTTP 400

MOBILE

Closing Meeting — findings list, agreed timeframes, the accuracy decision, and two on-device signature pads.

STEP 09 One signature does four things

REPORT · Auditor · HARD STOP

"The report cannot be issued without the auditor's signature. Signing triggers distribution and creates the corrective actions."

Signing is not a formality on top of a transition — it is the transition. When the auditor signed AUD-000010, four things happened in one movement:

- the report was generated from the data — scores, findings, benchmark, clause mapping
- a corrective action was created for every non-conformance, with the deadlines agreed at the closing meeting
- the report was distributed to the supervisor and the Safety Manager
- the site's history was re-evaluated for the persistent-poor-performance rule

The report the system built, with no document written by hand:

Element	AUD-000010
---------	------------

Reference	RPT-000010
Score	62.5% — Acceptable
Overall rating	Unsatisfactory
Findings	5 (2 conformance · 1 observation · 1 minor NC · 1 major NC)
Benchmark	–20.5 points against AUD-000009's 83%
Repeat findings	0
Clause mapping	ISO 45001 8.2 — 3 findings, worst Major NC · OSHA 1910.157 — 1 finding
Corrective actions	CAPA-000236 (due 26 Aug) · CAPA-000237 (due 20 Sep)
Signed by	Auditor One · factual accuracy confirmed by Supervisor One

There is no way to have one without the other.

No path issues a report without raising the corrective actions, and none raises actions from a report nobody signed. Tying them to the same act is what stops findings being reported and then quietly forgotten.

Then it moves to the desktop

The full report — a long document with sections, findings, benchmark comparison and clause mapping — is reviewed and distributed from the web console, where a long document is genuinely easier to work with. Two gates apply, in order:

Who	What they do	Gate
Safety Manager	Reviews and approves the report before wider distribution	Cannot happen before the report is signed
Admin	Owens distribution beyond the site	Refused until the Safety Manager has approved

That second gate was tested: distributing before approval returns “The Safety Manager reviews and approves the report before wider distribution”. Without it, the review would be decorative.

MOBILE Audit Report — the summary, the signature pad, and Sign & Issue.

WEB Report Review — the full document, approve for distribution, release beyond the site, print.

STEP 10 The audit stays open

CLOSE · The System + Auditor · HARD STOP

AUD-000010's report was issued. The audit did not close. It moved to step 10 with two open findings and stayed there.

"Keeps the audit open until every action is verified. Confirms findings were genuinely closed, not just marked closed."

Marking complete is a claim; verification is a visit

The auditor returns at 30, 60 and 90 days and checks each fix on site. The system will not let them verify a finding whose corrective action is still open:

"Corrective action CAPA-000236 is still Open. It has to be completed before the finding can be verified." — the system's actual refusal, HTTP 400

Once both actions were completed and both findings verified effective, the audit closed — status completed, all ten steps done, stage 8 CLOSE.

And if the fix is not holding

A failed effectiveness check does not just reopen the finding. It reopens the corrective action, and it reopens the audit itself — even one already closed. Tested on AUD-000010: a failed 60-day check took a completed audit back to capa_open, step 10 blocked, stage 5 IMPROVE.

This was a real bug, found and fixed.

The first implementation treated an audit as closed if its status said completed, which meant clearing the closure date did nothing — the status the system was about to recompute answered "still closed" and overwrote the reopen with itself. A migration backfilled closure dates onto historic audits so a single field is the authority.

MOBILE

Track Findings Out — per-finding verification with GPS, and the cross-audit 30/60/90 queue.

WEB

Trends & Oversight — cross-site comparison, repeat-finding analysis, and the re-audit decision.

The five rules that fire on their own

Each fires without anyone needing to notice.

None of these is a screen somebody has to open or a report somebody has to run. Each is raised by the transition that creates the condition, at the moment it is created.

Trigger	What happens	Fires at
Critical finding on site	Safety Manager and executive notified immediately. Work may be suspended before the audit finishes.	The moment the item is answered
Major non-conformance	Safety Manager notified within 24 hours. A corrective action must exist within 7 days.	Classification, step 7
Regulatory finding	Triggers the statutory notification workflow with its own legal deadline.	Classification, step 7
Audit not conducted	Alert at 110% of the scheduled date. A missed audit is itself a finding.	Daily sweep

Persistent poor performance	Two Major NCs at one site in 12 months, or below 65% twice running → mandatory re-audit within 30 days.	Report issue, step 9
------------------------------------	---	----------------------

The re-audit decision belongs to a person

The trigger fires automatically. What to do about it does not. The Safety Manager either schedules the re-audit — which creates it, seeds its checklist and links it to the original — or waives it. Waiving requires a written reason, and the system refuses without one.

An unexplained waiver of a mandatory re-audit is the single most useful thing for a regulator to find, which is precisely why it cannot be recorded silently.

Who does what

The specification's job titles, mapped to this platform's roles.

The source documents name several audit roles — Lead Auditor, Audit Team, Auditee Management, Safety Advisor, ISMS Director. Those are job titles in a large safety department, not app roles. On this platform they resolve to the four mobile roles plus web admin.

Document role	Platform role	Note
Lead Auditor + Audit Team	AUDITOR (mobile)	One role; where several are assigned to one audit, one is designated lead
Auditee Management	SUPERVISOR (mobile)	The supervisor of the area under audit
Workers interviewed	WORKER (mobile)	Observed and questioned during the walk
Safety Manager / Safety Advisor	SAFETY MANAGER (mobile + web)	Authorises, assigns, approves
ISMS Director / Safety Director	ADMIN (web only)	No mobile app; oversight and programme ownership only

In one line each

AUDITOR	Conducts the field inspection · collects and links evidence · classifies every finding · signs the report — nothing issues without it · verifies effectiveness at 30/60/90
WORKER	Interviewed and observed · shows competence and records · completes actions assigned to them · receives the toolbox talk
SUPERVISOR	Gets two weeks' notice · attends opening and closing · confirms factual accuracy · owns the actions for their area
SAFETY MANAGER	Authorises the programme · assigns an independent auditor · approves the report · owns the re-audit decision
ADMIN	Owens the audit programme · maintains checklist templates · receives every report · owns cross-site trends
THE SYSTEM	Builds the schedule and brief · calculates scores and thresholds · raises actions from findings · fires the re-audit trigger

Why the work splits where it does

“Everything requiring observation, evidence or a signature happens where the work is. Everything requiring reading, comparison or distribution happens where the screen is bigger. The auditor never re-enters anything they already recorded in the field.”

Step	Where	Screen
01 Schedule generated	Web	Audit Programme
02 Team assigned	Web	Audit Register
03 Pre-audit prep	Mobile	Brief Pack
04 Opening meeting	Mobile	Opening Meeting
05 Field inspection	Mobile	Field Inspection
06 Evidence captured	Mobile	Field Inspection · Worker Interview
07 Findings & score	Mobile	Findings & Score
08 Closing meeting	Mobile	Closing Meeting
09 Report issued	Mobile signs · Web reviews	Audit Report · Report Review
10 Findings tracked out	Mobile verifies · Web oversees	Track Findings Out · Trends

What the phone has to be able to do

Capability	Why	Status
Offline execution	Audits happen where there is no signal. Not optional for a field tool.	Built
Camera & evidence linking	Photos attached to the checklist item, not a folder.	Built
GPS stamping	Proves the observation was made where claimed.	Built
On-device signature	Both parties sign before leaving site, so findings lock immediately.	Built
Voice-to-text	Hands are often gloved or occupied.	OS dictation
QR / barcode scan	Identify equipment without typing a reference.	Typed entry — needs a scanner library

What was built

For the record — the shape of the implementation.

Data

Table	Holds
audits	The audit and its whole lifecycle — trigger, band, meetings, signatures, score, report
audit_checklist_items	The checklist as rows, so a section can score on its own
audit_findings	One row per classified finding, tracked out individually
audit_evidence	Photos, documents, scans, notes and worker interviews, bound to a checklist line
audit_programme	Each site's band, cadence, authorisation and generation history
audit_checklist_templates	The versioned templates every audit runs from

Judgement, kept out of the controllers

Module	Owns
audit_scoring	The rubric, the bands, the five classifications, the overall rating
audit_programme	The six triggers, the four bands, the frequency table
audit_steps	The ten steps, their owners, their hard stops, and the status they imply
audit_escalation	The five triggers that fire on their own
audit_brief	The step 03 brief pack
audit_calendar	Calendar generation and the 14-day reminder
audit_trends	Cross-site comparison and repeat-finding analysis
audit_templates	Template resolution and versioning

Screens

Mobile (Auditor)	Web console
Dashboard — queue, programme, escalations	Audit Programme — bands, authorisation, generation
Assigned Audits — step state per audit	Audit Register — the queue and assignment
Audit Detail — the one action that is due	Report Review — the long document, approve, distribute
Brief Pack · Opening Meeting	Trends & Oversight — cross-site, repeats, re-audit
Field Inspection · Worker Interview	Templates & Auditors — versioned editor, register
Findings & Score · Closing Meeting	
Audit Report · Track Findings Out	

Honest limits

Three things are worth stating plainly rather than leaving to be discovered:

- **No scanner.** QR and barcode scanning is typed entry. No scanner library is installed, and shipping a button that pretends to scan would be worse than one that plainly asks for the reference.
- **Dictation, not integration.** Voice-to-text relies on the phone keyboard's dictation key rather than an in-app speech integration.
- **Not device-tested.** The mobile app is type-checked and the web app builds, but neither has been walked through on a device against live data. The backend flow has — every number in this document came from that run.

The auditor is one of four mobile roles. Steps 4 to 8 are conducted in the field on the phone, offline-capable, with evidence and signatures captured on site. An audit is not closed when the report is

issued — it stays open until every corrective action it raised has been verified effective.